

USENIX Security '25 Artifact Appendix: Relocate-Vote: Using Sparsity Information to Exploit Ciphertext Side-Channels

Yuqin Yan[†], Wei Huang^{†‡}, Ilya Grishchenko[†], Gururaj Saileshwar[†], Aastha Mehta^{*}, and David Lie[†]

[†]University of Toronto, [‡]Seneca Polytechnic, ^{*}University of British Columbia
 yuqin.yan@mail.utoronto.ca, wei.huang1@senecapolytechnic.ca, gururaj@cs.toronto.edu,
 aasthakm@cs.ubc.ca, {ilya.grishchenko,david.lie}@utoronto.ca

A Artifact Appendix

A.1 Abstract

This artifact includes the source code, scripts, and datasets used to support the main claims, figures, and tables presented in the paper. The paper demonstrates how a malicious hypervisor can infer sparsity patterns in the memory of victim confidential Virtual Machines (CVMs) protected by AMD SEV-SNP. This is accomplished by identifying encrypted memory blocks that contain either prevalent or non-prevalent values, exploiting ciphertext side channels via the `SNP_PAGE_MOVE` command—designed initially for resource management.

At a high level, the artifact provides: (1) a malicious hypervisor implemented as part of a modified Linux kernel, (2) code simulating the behavior of a victim CVM, (3) monitoring logic used by the hypervisor to observe the victim through controlled channels and ciphertext side channels, and (4) analysis tools to recover and evaluate sparsity information from the victim’s memory.

A.2 Description & Requirements

A.2.1 Security, privacy, and ethical concerns

This artifact is based on a modified kernel implementation, SEV-Step. Both the original version and our implementation prioritize the functionality of attack primitives over kernel stability. As a result, the system may experience unpredictable behavior during evaluation, including kernel crashes, instability, and unresponsiveness. **WARNING:** This artifact is intended strictly for research and evaluation purposes. **Do not use it in production environments.**

Known limitations and usage guidelines:

- Do not run multiple VMs simultaneously.
- Do not assign multiple cores to a single VM, which may prevent the VM from booting.
- Avoid forcefully terminating an evaluation program, which may cause system issues.

- The CVMs cannot be correctly shut down sometimes. Specifically, you cannot terminate a running CVM by either Ctrl-C in the terminal with the “login” prompt or `pkill`. If this happens, you need to reboot the host machine.

A.2.2 How to access

The artifact has been uploaded to Zenodo at <https://doi.org/10.5281/zenodo.15609905>.

A.2.3 Hardware dependencies

The implementation is based on SEV-SNP, and the proposed primitive leverages the `SNP_PAGE_MOVE` relocation command. Notably, AMD fixed a relevant implementation bug in firmware version 1.55 (build 20). Therefore, for efficient page relocation, firmware version 1.55.20 or later is required.

A.2.4 Software dependencies

Toolchain in the host platform: The host platform is running Ubuntu 24.04.1 LTS. For installed packages, check the section “Evaluation Platform” in `README.md` for details.

Environments supporting `pip install`. Examples of such environments are Anaconda and `python-venv`. Recently, `pip install` in the native, raw environment is prohibited due to PEP 668—Marking Python base environments as “externally managed”.

Blender. The software Blender is required to assemble the pieces leaked from OpenVDB operations manually. The installation instructions can be found at <https://www.blender.org/download/>.

A.2.5 Benchmarks

Dataset. We rely on the external datasets, CQ500. Although it is externally available¹, considering the permanent ac-

¹<https://www.kaggle.com/datasets/crawford/quareai-headct>

cessibility, we included necessary components in our artifact package, including a subset of the artifact located in `openvdb-leak/data/quareai-headct`. The distribution is based on the original license, CC BY-NC-SA 4.0, and the license file `LICENSE` and corresponding `README.md` are in the directory.

Weights of large language model (LLM). Our experiments relied on the weights of `ReLuLLaMA-13b`² being served by PowerInfer. For the ease of reproducibility, we include the weights in the artifact package. The weights are released under the llama2 license. We have included license files `LICENSE.txt` and `Notice`.

A.3 Set-up

A.3.1 Installation

The installation consists of four steps. Follow the detailed instructions in the “Installation” section of `README.md`.

1. Set `RV_ROOT_DIR` to the root directory of the artifact by executing the script in it:

```
source ./env.sh
```

2. Create the CVM images, `aslr.qcow2` and `big-disk.qcow2` (Table 1) in `$RV_ROOT_DIR/vm-images-creation/`. The first line takes around 10 minutes to execute. The second line can take around 30 minutes. The default username is `ubuntu` and the password is `123456`.

```
$RV_ROOT_DIR/vm-images-creation/create_vm1.sh
$RV_ROOT_DIR/vm-images-creation/create_vm2.sh
```

3. Build and install the customized Linux kernel.

```
cd $RV_ROOT_DIR/sev-step
./build.sh ovmf && \
./build.sh qemu && \
./build.sh kernel
```

The above commands should generate `kernel-packages/` in `$RV_ROOT_DIR/sev-step/`. Install the `*.deb` packages using the following command:

```
sudo dpkg -i \
$RV_ROOT_DIR/sev-step/kernel-packages/*.deb
```

After installation, execute `sudo update-grub` to see if the following lines are in the output:

```
Found linux image: /boot/vmlinuz-5.19.0-rc6-sev-step-999ae99
Found initrd image: /boot/initrd.img-5.19.0-rc6-sev-step-999ae99
```

If so, the installation should be successful. Record its index in the output, where the index starts from 0. For example, if the output is:

```
Generating grub configuration file ...
Found linux image: /boot/vmlinuz-6.8.0-49-generic
Found initrd image: /boot/initrd.img-6.8.0-49-generic
Found linux image: /boot/vmlinuz-5.19.0-rc6-sev-step-999ae99
Found initrd image: /boot/initrd.img-5.19.0-rc6-sev-step-999ae99
...
```

The index of the customized kernel is 2 (the third line, starting from 0). Now update the `GRUB_DEFAULT` in `/etc/default/grub` to boot into the new kernel by default.

```
GRUB_DEFAULT="1>[index]"
```

Here, `[index]` is the index of the new kernel in the GRUB menu just recorded. Then update GRUB configuration:

```
sudo update-grub
```

Reboot the machine and verify the kernel version with `uname -r`. It should output `5.19.0-rc6-sev-step-999ae99`.

4. Compile the user-space projects for the experiments.

```
cd $RV_ROOT_DIR && make
```

A.3.2 Basic Test

Launch guest CVMs. Run to launch **VM1** (Table 1):

```
$RV_ROOT_DIR/launch-vm1.sh
```

Upon a successful launch, a login prompt should appear in the terminal (pressing “Enter” is required sometimes). If the launch fails, check the script’s output for potential troubleshooting instructions. Rebooting the host machine to reset the state if necessary.

For launching **VM2**, execute the following instruction in a new terminal (`source ./env.sh` first in it), which also kills **VM1** first due to the implementation limitation that **VM1** and **VM2** cannot be run simultaneously. When switching VMs, read the script output carefully for potential instructions to reboot the host machine if necessary.

```
$RV_ROOT_DIR/launch-vm2.sh
```

VM Name	qcow2 image	memory-size (in MB)
VM1	<code>aslr.qcow2</code>	4096
VM2	<code>big-disk.qcow2</code>	4096

Table 1: CVMs used in the evaluation.

Test attack primitive. After launching a CVM, execute the following command in a different terminal than the one with the login prompt:

²<https://huggingface.co/PowerInfer/ReluLLaMA-13B-PowerInfer-GGUF>

```
$RV_ROOT_DIR/t1_test-primitive.sh
```

If it succeeds, a green checkmark followed by “Successfully...” will appear on the screen when the script finishes.

A.4 Evaluation workflow

A.4.1 Major Claims

- (C1):** The frequency distributions of plaintexts in the CVM are preserved when encrypted at the same system physical address (sPA), which can be exploited to learn the ciphertexts corresponding to the prevalent values in the CVM, such as zero (Section 1, Section 3.1).
- (C2):** The sparsity information in the guest page tables can be exploited to reduce significantly the number of possible guest virtual addresses (GVA) of the `glibc` library (Section 4, exemplified on `nginx`).
- (C3):** The sparsity information in the OpenVDB node buffers can be used to leak the structural information of the objects being processed under OpenVDB library operations (Section 5).
- (C4):** The sparsity information in the ReLU buffers can be decoded and utilized to leak information about tokens by placing probes on the blurred activation information (Section 6).

A.4.2 Experiments

Claim C1 is supported by experiment **E1-C1**. For claim C2, the main supporting experiment is **E2-C2**, plus an optional experiment E3-C2. In contrast to the main experiments, which take minutes, experiments marked as (optional) require significant time (several hours) to finish. Although we believe they are not required to support the claims, we provide them for completeness. We support claim C3 with the main experiment **E4-C3** followed by optional experiments E5-C3, E6-C3, and E7-C3. Finally, we support claim C4 with the experiment **E9-C4** preceded by the optional experiment E8-C4.

(E1-C1): [learning-zero-ciphertexts] [1 compute-minutes + 3 human minutes] In this experiment, the attacker learns the ciphertexts corresponding to plaintext zeros in the CVM on the target page.

Preparation: (1) Launch a victim CVM:

```
$RV_ROOT_DIR/launch-vm1.sh # VM1 as an example
```

Execution: Once the CVM is launched, in another window (or shell, log-in session)³, execute

```
$RV_ROOT_DIR/e1_ciphertext-learning.sh
```

³Depending on your preference, you can achieve multiple shells by using `tmux` or logging in to a new session.

Results: If it succeeds, a green check followed by “Successfully...” will appear on the screen when the script finishes. A detailed description of the contents of collected logs, including `$RV_ROOT_DIR/cl.log` and `$RV_ROOT_DIR/dmesg-output.log` is in `README.md`.

(E2-C2): [ASLR] [Around 6 compute-minutes + 5 human-minutes]: This experiment is a minimal demonstration of the ASLR de-randomization scenario, exemplified by de-randomizing the `glibc`’s base address of `nginx` with a single memory layout.

Preparation: (1) Launch the victim VM1 for ASLR.

```
$RV_ROOT_DIR/launch-vm1.sh
```

Execution: Execute

```
$RV_ROOT_DIR/e2_aslr.sh
```

Results: The output logs are located in `$RV_ROOT_DIR/aslr/aslr-min-log/`. The detailed structure of this directory is provided in the "Description of the Output Logs" section of the `README.md` file. Upon successful execution, the directory `0/nginx-worker/` should contain a file whose name ends with `window_size_64.log`. The results of non-windowed and windowed trackings can be checked at the end of `0/nginx-worker/windowed_evaluation.log`, starting with the line “ASR metrics for service `nginx-worker`”.

(E3-C2): [ASLR-full (optional)] [5-6 compute-hours + 10 human-minutes] The spirit of this experiment is the same as E2, but extends to all five applications in Section 4 and evaluates 10 layouts for each. As a result, it takes around 50× more time than it takes for the **E2-C2** experiment.

Preparation: The same as **E2-C2**.

Execution: Execute

```
$RV_ROOT_DIR/e3_aslr-full.sh
```

Results: The output logs will be in `$RV_ROOT_DIR/aslr/aslr-all-log/`. They are similar to **E2-C2** but with more memory layouts and more applications for each layout. Upon completion, execute the following scripts to parse the logs and produce data for Table 2 in the paper.

```
cd $RV_ROOT_DIR/aslr
python ./scripts/aslr_track_analysis.py
```

(E4-C3): [openvdb-traversal] [10 compute-minutes, 15-25 human-minutes] This experiment leaks the object’s structural information from OpenVDB’s read-only traversal operation.

Preparation: (1) Launch the victim CVM VM2.

```
$RV_ROOT_DIR/launch-vm2.sh
```

Execution: Execute

```
$RV_ROOT_DIR/e4_openvdb-traversal.sh
```

Results: The scripts produce logs in `$RV_ROOT_DIR/openvdb-leak/openvdb-traversal-1-time/`. A detailed description of its structure is in `README.md`. In this directory, the file `output_all.txt` contains the page fault events during tracking. The scripts process the record of page-fault events and output PLY files (a 3D data format commonly used to store point clouds and polygonal meshes). When the recovery succeeds, it should contain eight PLY files in `evaluation_traversal_extraction_<timestamp>_1/<ply-file-dir>/`.

Optional: Visualize the extracted pieces and assemble them. Use `scp` or `rsync` the `<ply-file-dir>` to local machine. We also provide a Blender project at `$RV_ROOT_DIR/openvdb-leak/blender_projects/traversal.blend`, which you can transfer to your local machine using `rsync` and open with Blender. The object labeled “original” in the upper-right corner corresponds to the source object shown in Figure 8(a) of the paper, where the source file is provided as `$RV_ROOT_DIR/openvdb-leak/blender_projects/original.ply`.

After copying the PLY files to your machine using `rsync`, import them into Blender. Manually reassemble the eight fragments by adjusting their positions along the X, Y, and Z axes, as illustrated in `$RV_ROOT_DIR/openvdb-leak/blender_projects/blender-adjust-xyz.png`. Once completed, make only the assembled pieces visible by toggling the eye icons in the top-right panel. Finally, export the visible parts as a single PLY file (`traversal.ply`).

We provided some analysis tools in `$RV_ROOT_DIR/openvdb-leak/blender_projects`. For example, you can use `visualize_traversal.py` to visualize the generated `traversal.ply`, achieving a visualization effect similar to that in Figure 8 of the paper. Additionally, `p2p_dist_traversal.py` generates part of the comparison in Figure 9 of the paper, contrasting the distribution of nearest-neighbor distances between objects extracted from the traversal operation and randomly populated points.

(E5-C3): [openvdb-traversal-3-trackings (optional)] [Triple the time required for E4 + around 6 GB disk space for storing the logs.] Everything remains the same as in **E4-C3**, except the experiment is repeated three times for evaluation.

Preparation: The same as **E4-C3**.

Execution: The same as **E4-C3**’s execution, but pass a different number as the number of times to run as a parameter. Execute:

```
$RV_ROOT_DIR/e5_openvdb-traversal-3.sh
```

Results: The same as **E4-C3** but applied to more tracking instances.

(E6-C3): [openvdb-construction (optional)] [Around 3

compute-hours + 15-25 human-minutes + around 80GB disk space for storing logs] This experiment leaks the object’s structural information from OpenVDB’s construction operation.

Preparation: The same as **E4-C3**.

Execution: Execute

```
$RV_ROOT_DIR/e6_openvdb-construction.sh
```

Results: The same to **E4-C3** but applied to a different victim’s operation.

(E7-C3): [openvdb-construction-3-trackings (optional)] [Triple the time required for E6-C3 + around 240 GB disk space for storing the logs.] Everything remains the same as in **E6-C3**, except the experiment is repeated three times.

Preparation: The same as **E4-C3**.

Execution: Execute

```
$RV_ROOT_DIR/e7_openvdb-construction-3.sh
```

Results: The same as **E6-C3** but applied to more tracking instances.

(E8-C4): [sparse-llm-data] [5–6 compute-hours + 3 human-minutes + around 230GB disk.] This experiment produces data at `$RV_ROOT_DIR/sparsellm-probe/powerinfer_reduced_activation_probes_results/` for **E9-C4**, which generates Figures 11 and 12 in the paper.

Execution: Execute

```
$RV_ROOT_DIR/e8_sparse-llm-data.sh
```

The above scripts take 5–6 hours to execute.

Results: The result directory `$RV_ROOT_DIR/sparsellm-probe/powerinfer_reduced_activation_probes_results/` is generated and can be visualized in **E9-C4**.

(E9-C4): [sparse-llm-figures] [1 compute-minute + 3 human-minutes] Visualize the results of **E8-C4**.

Preparation: Finish **E8-C4**.

Execution: On the evaluation platform, execute:

```
$RV_ROOT_DIR/e9_sparse-llm-figures.sh
```

Results: Figures 11 and 12 in the paper can be produced as `probe_results.png` and `world_large_reduce4.png` in the `sparsellm-probe` directory when the execution succeeds.

A.5 Version

Based on the LaTeX template for Artifact Evaluation V20231005. Submission, reviewing and badging methodology followed for the evaluation of this artifact can be found at <https://secartifacts.github.io/usenixsec2025/>.